

Zoya Qureshi

+91 75666-05934 | zoyaqureshi28103@gmail.com | LinkedIn | GitHub

Summary

AI Engineer specialising in Generative AI and LLM-powered systems, with experience building and deploying production-grade, multi-user applications. Skilled in designing scalable backend architectures, optimising LLM performance, and delivering end-to-end AI products. Seeking remote AI/GenAI roles to build high-impact intelligent systems.

Experience

AI Development Intern

Codebyte Solutions — May 2025 – Nov 2025

- Engineered real-time GenAI pipelines integrating ASR, NLP, and LLM components using locally hosted models (Mistral, Phi-3).
- Reduced LLM inference latency by **90%** through 4-bit quantization using **bitsandbytes**.
- Designed stateful conversational agents using **LangGraph** for multi-turn memory and context retention.
- Built modular, scalable pipelines for real-time processing and improved system responsiveness.

AI Automation Intern

VKAPS IT Solutions Pvt Ltd — Apr 2025 – May 2025

- Developed LLM-powered prototypes using **LangChain**, accelerating proof-of-concept delivery.
- Automated workflows using **n8n**, reducing manual processing and improving efficiency.

Projects

FluentScroll — AI for Communication Training

[Live Demo](#)

- Designed and deployed a production-grade AI SaaS platform using **Next.js**, **FastAPI**, and **Supabase**.
- Architected secure multi-user system with **JWT authentication** and Row-Level Security (RLS).
- Implemented LLM-based content generation using **Llama 3.1** with caching, reducing redundant API calls.
- Built 60 FPS teleprompter engine using **requestAnimationFrame**.
- Designed session analytics pipeline with deduplication and real-time tracking.

Real-Time Voice Assistant (Health-Tech)

[GitHub Repo](#)

- Built end-to-end voice assistant using **Phi-3** and **faster-whisper**.
- Achieved **sub-100ms** intent routing latency using lightweight ML models.
- Integrated **InfluxDB** for real-time health data processing.

RAG System Development

[GitHub Repo](#)

- Developed an end-to-end RAG pipeline using **LangChain**, **FAISS**, and **HuggingFace embeddings** for accurate knowledge retrieval.
- Integrated **LLaMA 3.1 (Groq API)** for low-latency, context-grounded response generation with hallucination control.
- Implemented **intent-based routing and semantic search** across multiple knowledge bases.

Skills

Generative AI & LLM: LangChain, LangGraph, RAG, Prompt Engineering, Fine-tuning, Quantization, Mistral, LLaMA, Phi-3, HuggingFace, FAISS

Backend & MLOps: FastAPI, Docker, Redis, InfluxDB, Supabase, AWS SageMaker, DynamoDB

Machine Learning: PyTorch, TensorFlow, Scikit-learn, NLP, ASR, Intent Classification

Programming: Python (Advanced), SQL, Git

Systems: REST APIs, Async Processing, Caching, Scalable Architecture

Education

B.Tech, Artificial Intelligence & Data Science

Chameli Devi Group of Institutions, Indore — 2021 – 2025